RL methods for cart pole swing up and stabilization

Bogdan Alexandrov, Vadim Shirokinskiy

Problem statement

We have an cart pole, the pendulum is looking down. The task is to lift the pendulum up and hold it.

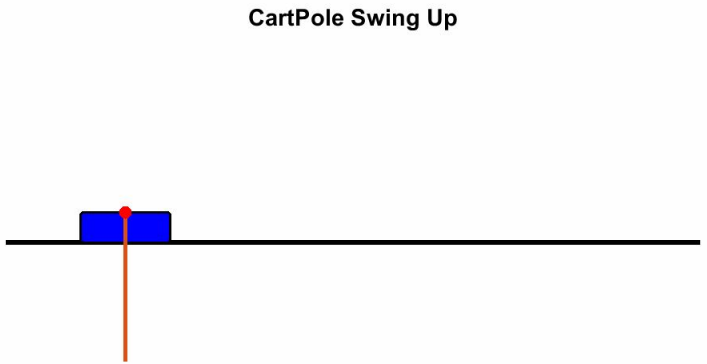

Parameters of the system

pendulum mass = 0.1

pendulum length = 0.5

area length = inf, 10, 5 (optionally)

steps per episode = 1500 - 3000

initial state: (pi,)

cart mass = 1

g = 9.81

dt = 0.003

System dynamic

$$\dot{\theta} = \omega$$

$$\dot{\omega} = \frac{(m + m_c)g \sin\theta - \cos\theta(u + ml\omega^2 \sin\theta)}{(4/3)(m + m_c)l - ml\cos\theta^2}$$

$$\dot{h} = dh$$

$$\dot{dh} = \frac{u + ml(\omega^2 \sin\theta - \dot{\omega}\cos\theta)}{m + m_c}$$

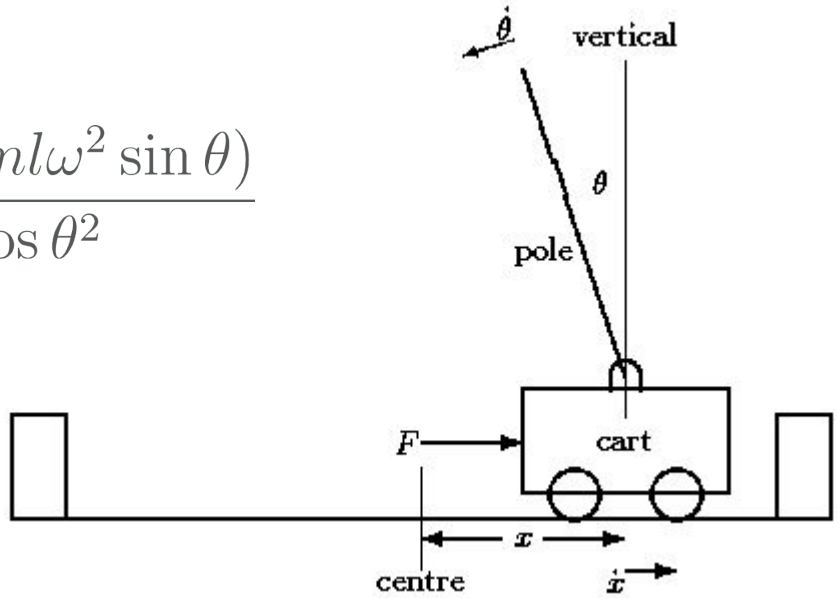

State -> observation transition

$$\theta, \omega \rightarrow \cos\theta, \sin\theta, \omega$$

Since we don’t care about absolute value of angle(10001 * pi or pi are equal for our system).

We need only sine and cosine of angle

Used approaches

RL methods:

1) Reinforce
2) Actor-Critic

Ways to influence the environment:

1) Apply continuous bounded action
2) Predict direction of predetermined force (used 10 in our experiments)

Setup1

Reinforce with fixed force. Cost is also fixed. Available surface is unlimited.

if $\cos\theta < -0.5$: return 1000

elif $\cos\theta < 0.8$:return 500

elif $\cos\theta < 0.95$: return 200

else: return 0

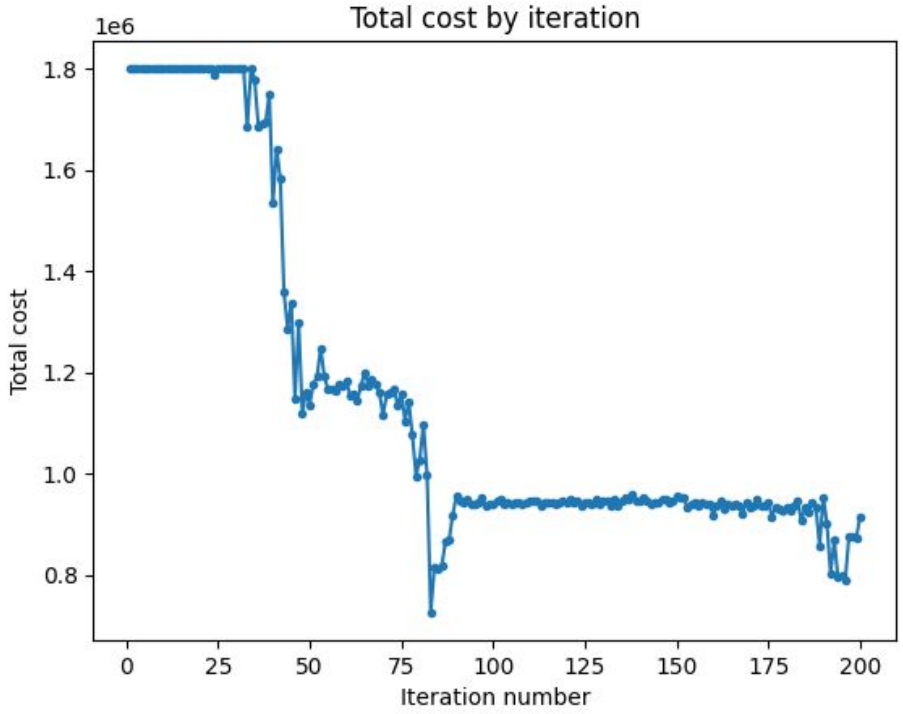

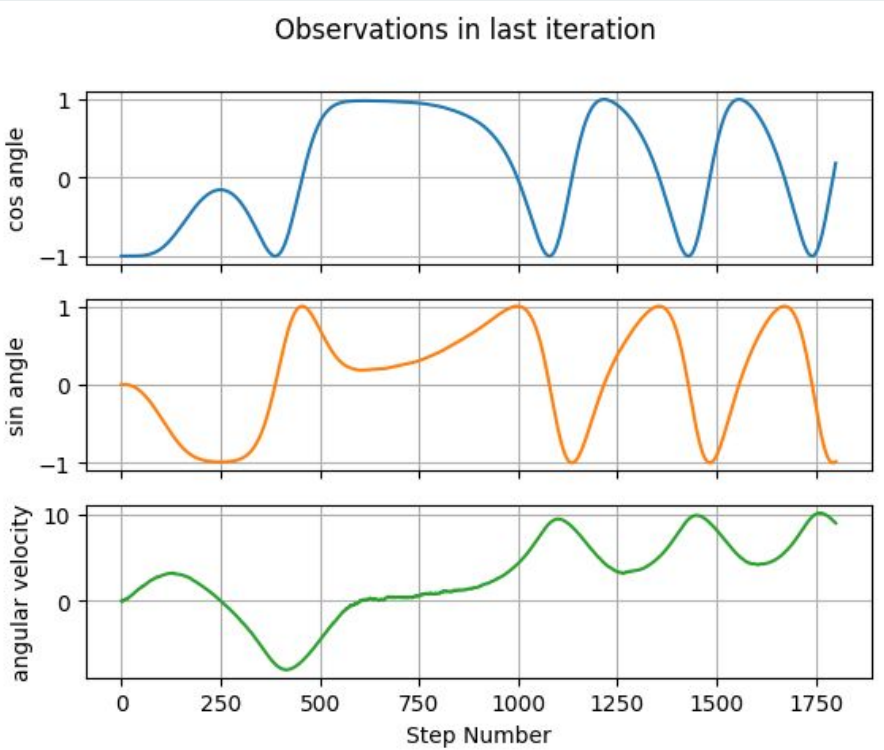

Hidden dim: 32
Hidden layers: 1

lr: 0.005
steps: 1800

N_episode: 2
N_iters: 200

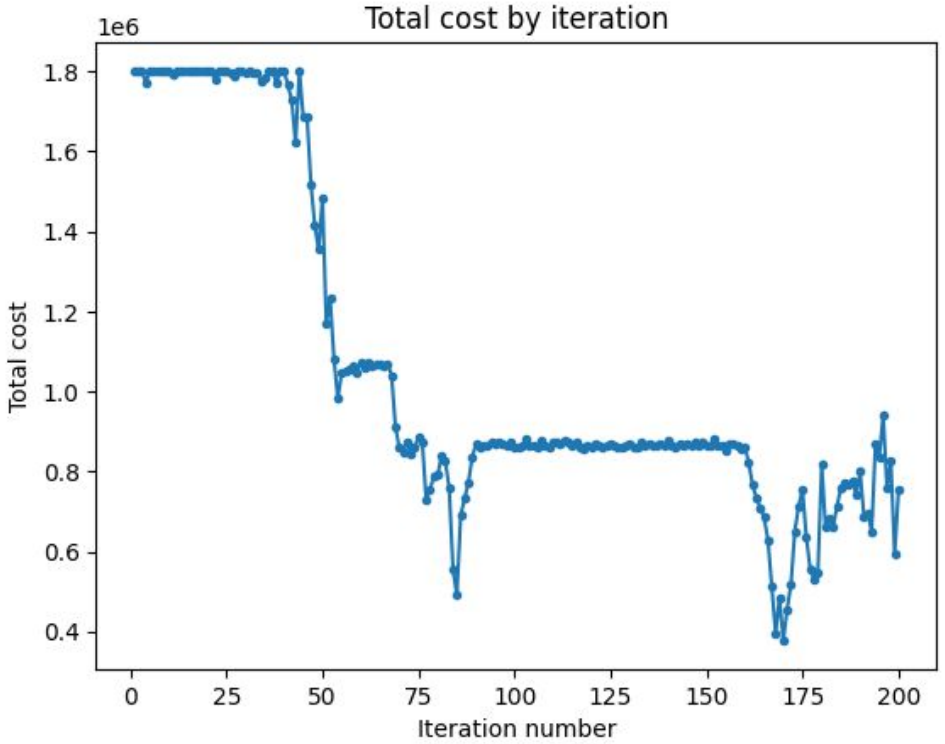

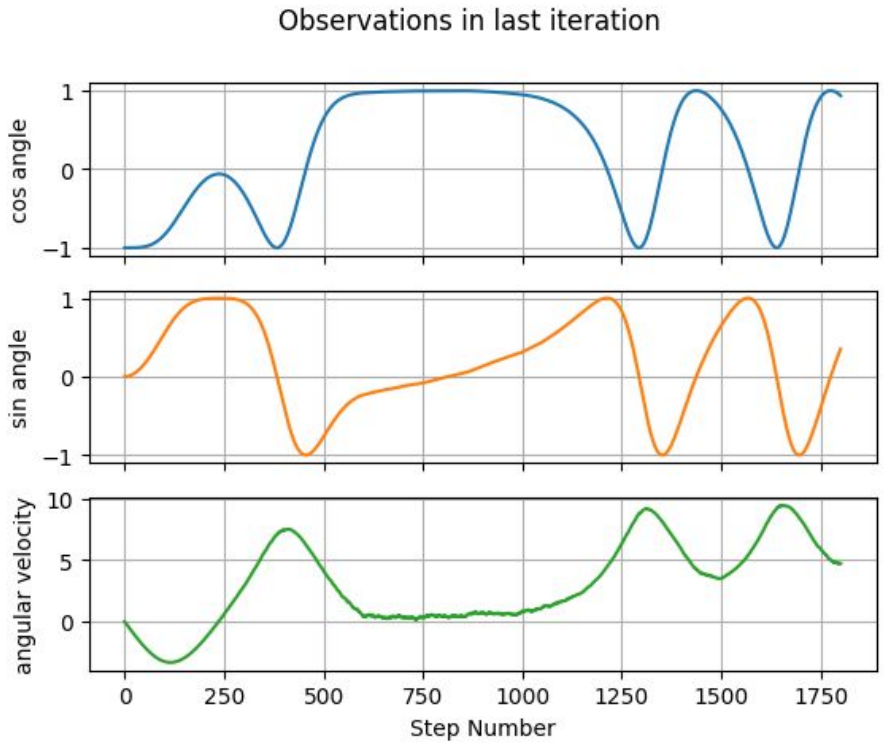

Hidden dim: 32
Hidden layers: 2

lr: 0.005
steps: 1800

N_episode: 2
N_iters: 200

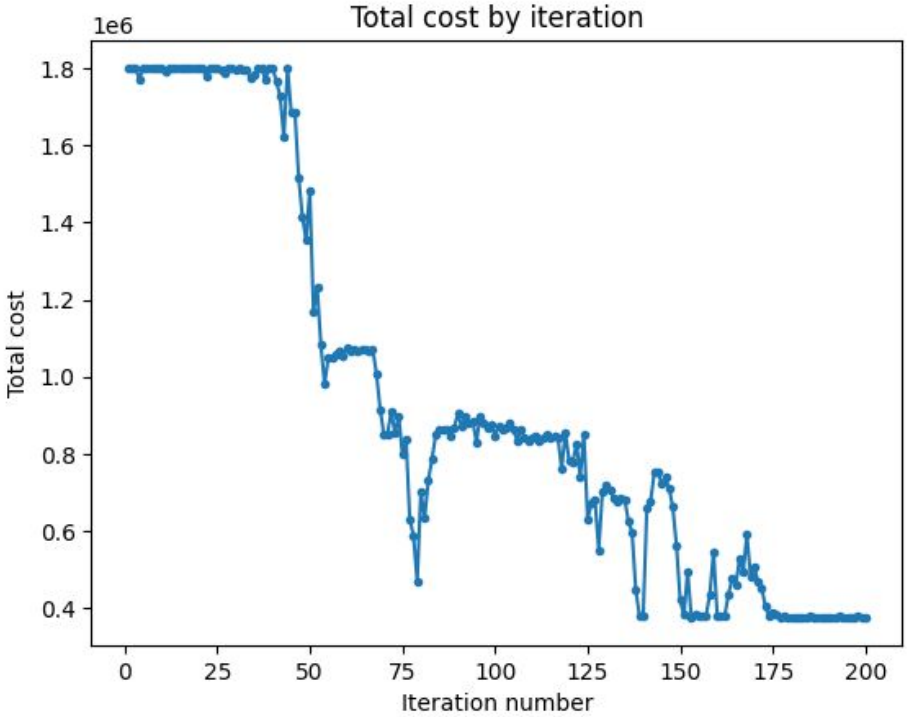

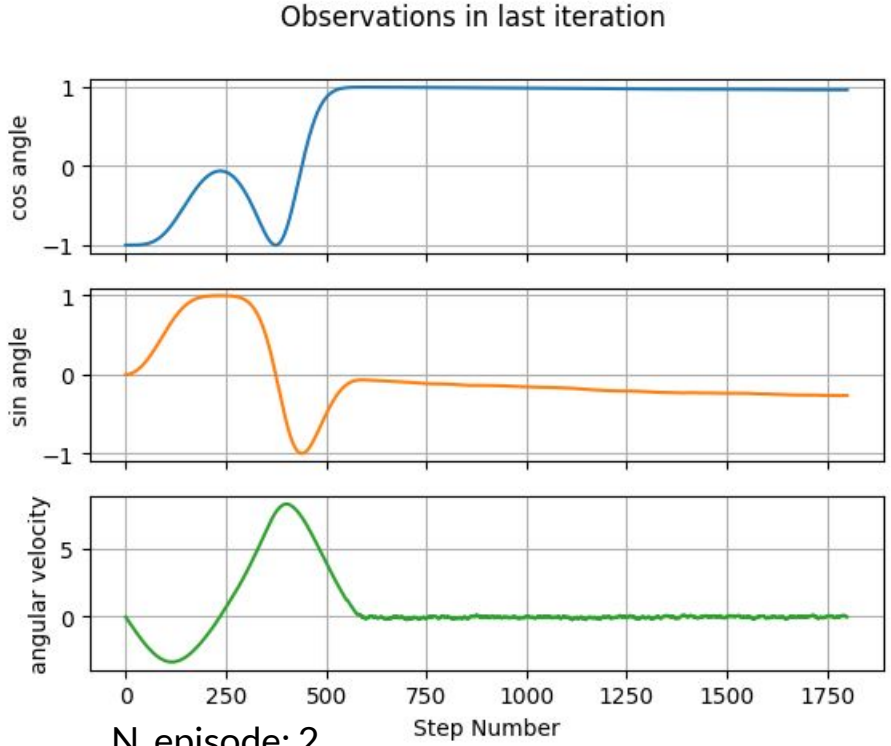

Hidden dim: 32
Hidden layers: 2

lr: 0.005
steps: 1800

N_episode: 2
N_iters: 200

Here added additional punishment on high velocity:

if $\cos\theta > 0.95 : \ \omega^2$

Setup 2

Same as previous, but changed cost to smooth one.

$$a * (1 - \cos\theta)^2 + b * \omega^2$$

where a, b are positive weights

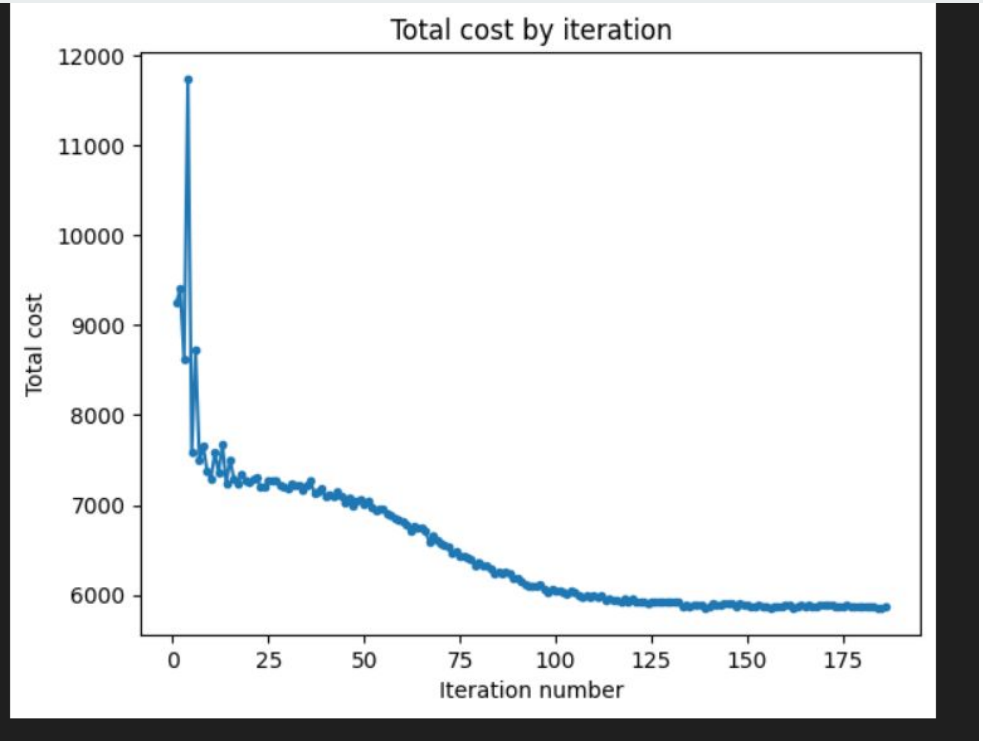

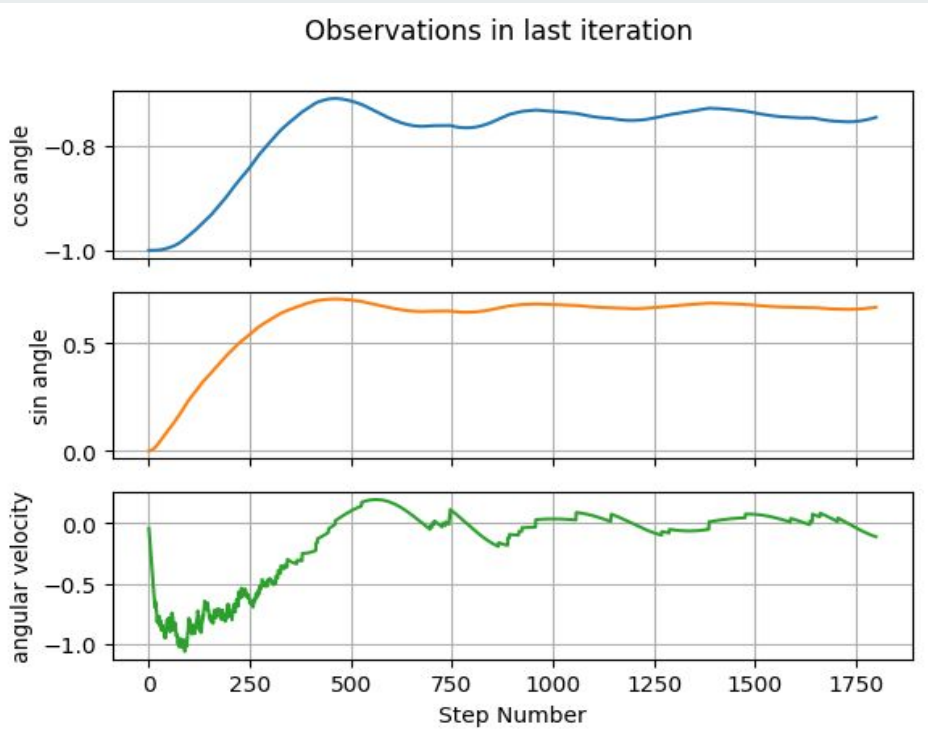

Hidden dim: 32
Hidden layers: 2

lr: 0.005
steps: 1800

N_episode: 2
N_iters: 200

cost coefs a=b=1.

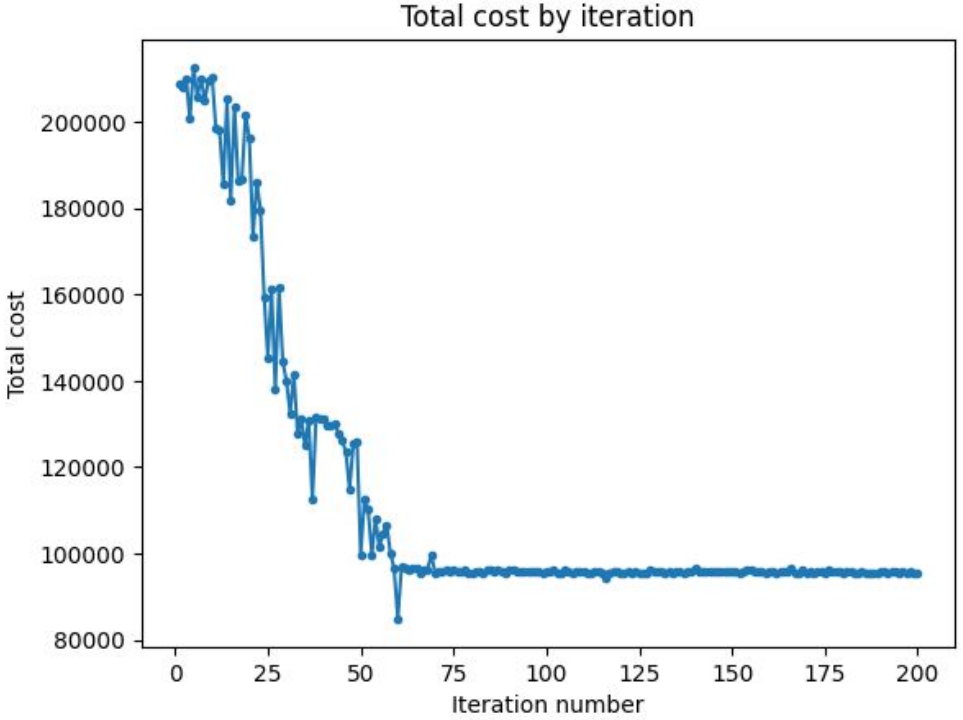

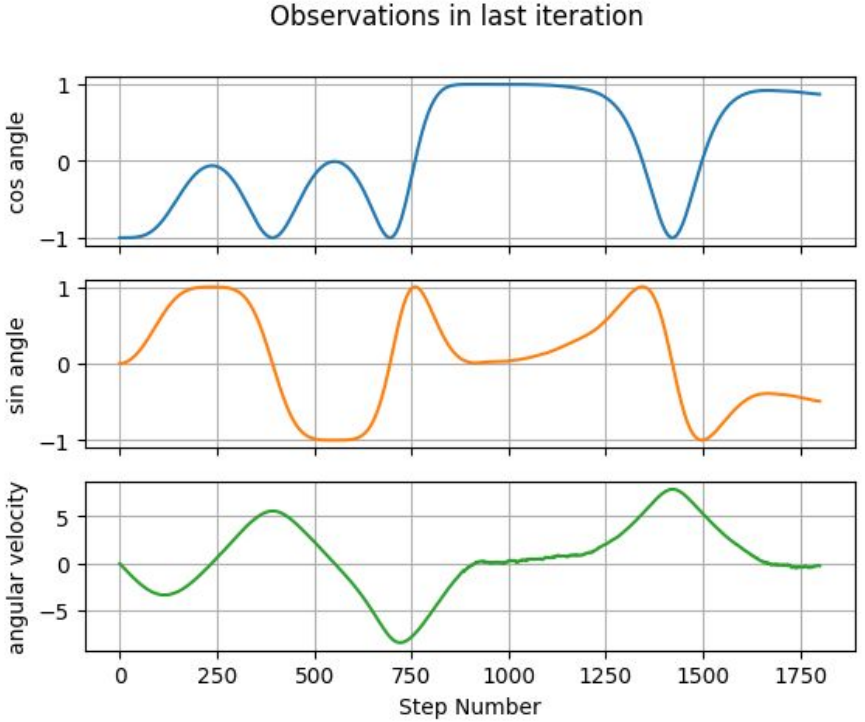

Hidden dim: 32
Hidden layers: 2

lr: 0.005
steps: 1800

N_episode: 2
N_iters: 200

cost coefs a=30, b=1.

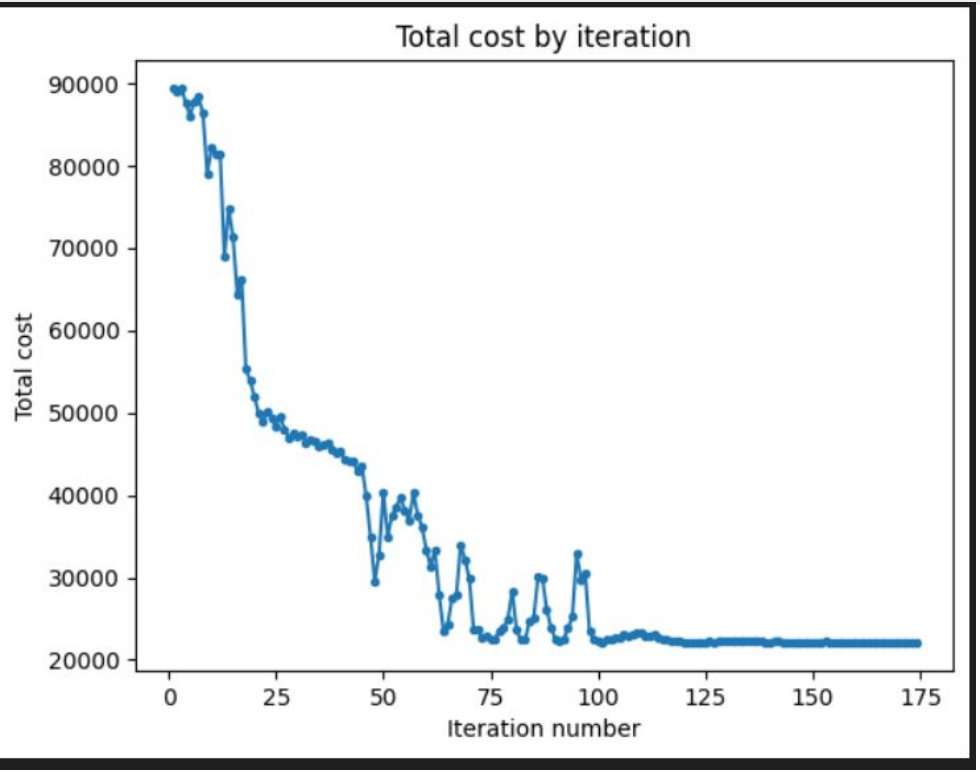

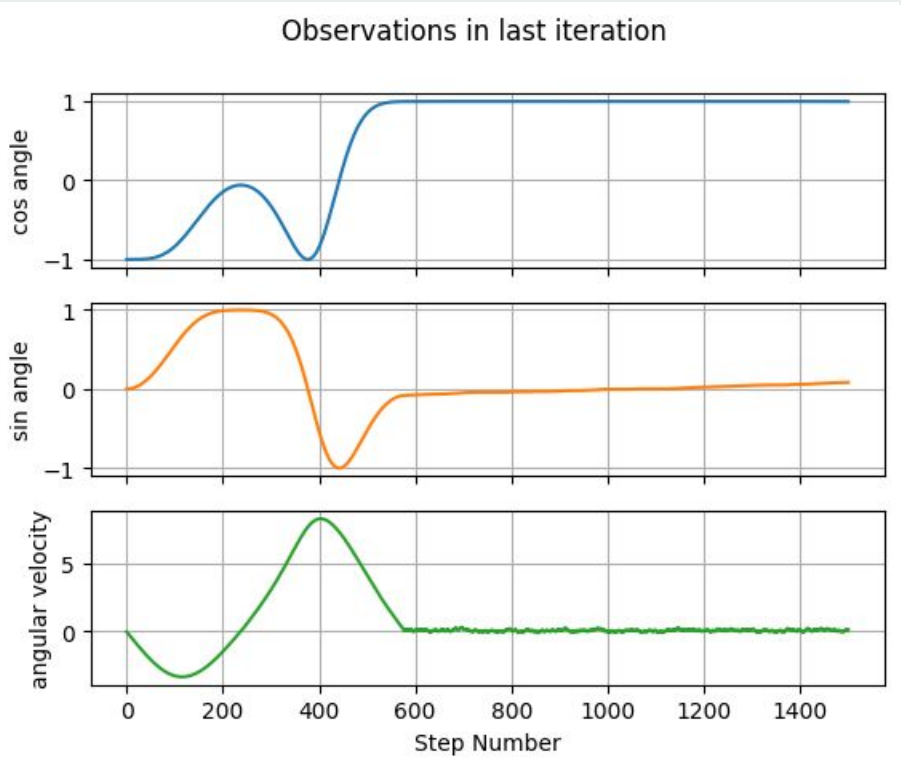

Hidden dim: 32
Hidden layers: 2

lr: 0.005
steps: 1500

N_episode: 2
N_iters: 200

cost coefs a=30, b=1.

Setup 3.

Here we changed model with fixed force to unfixed. Model predicts force itself within [-10, 10].

P.S. worked not good

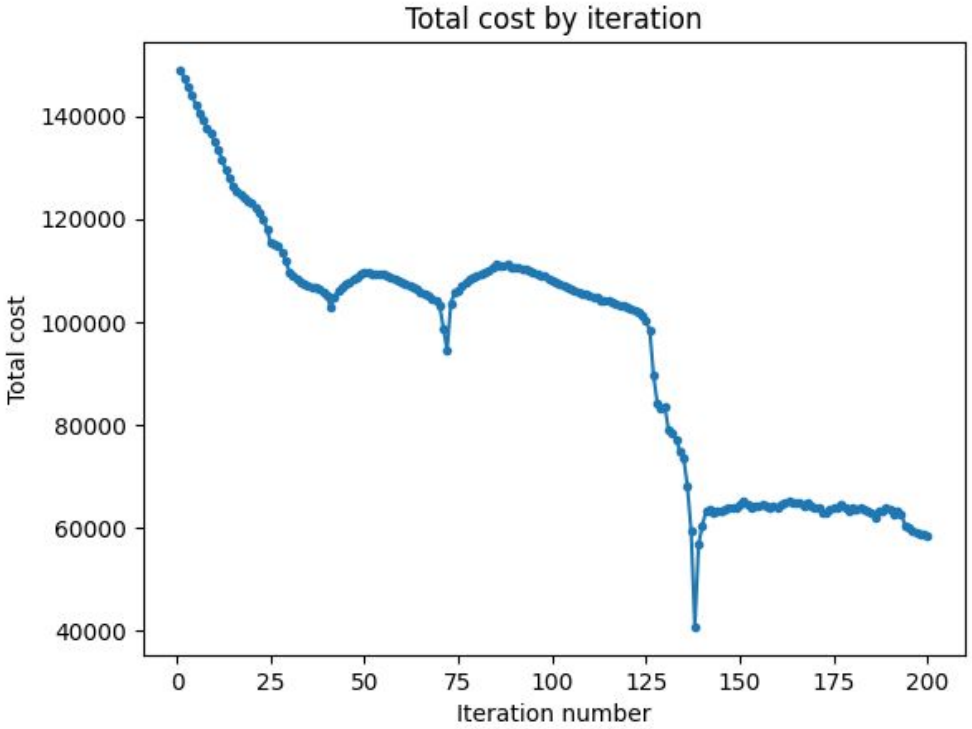

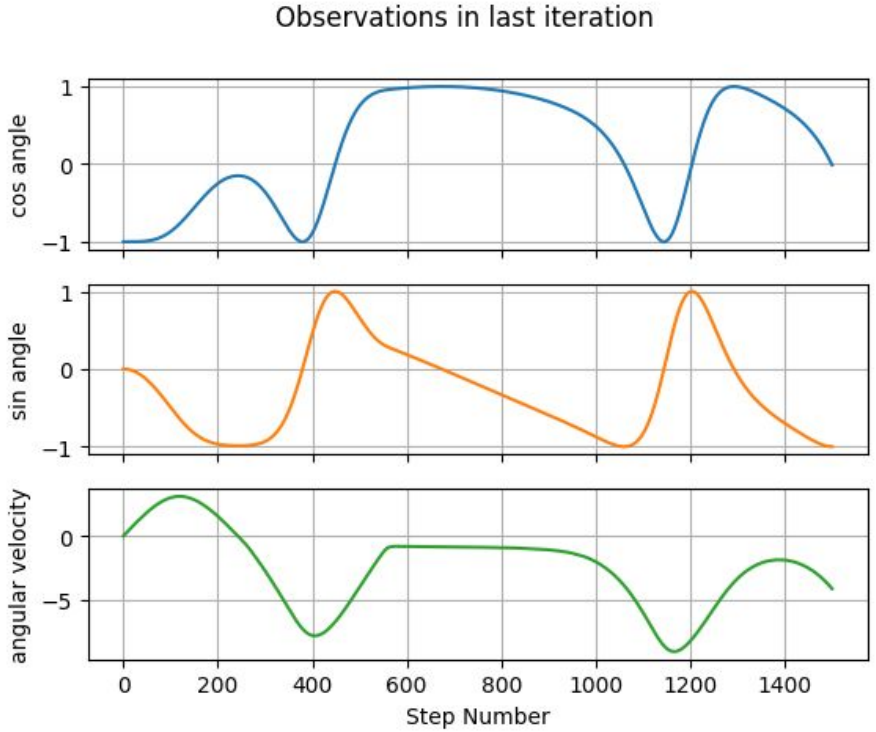

Hidden dim: 32
Hidden layers: 2

lr: 0.005
steps: 1500

N_episode: 2
N_iters: 200

cost function: $50 * (1 - \cos\theta)$

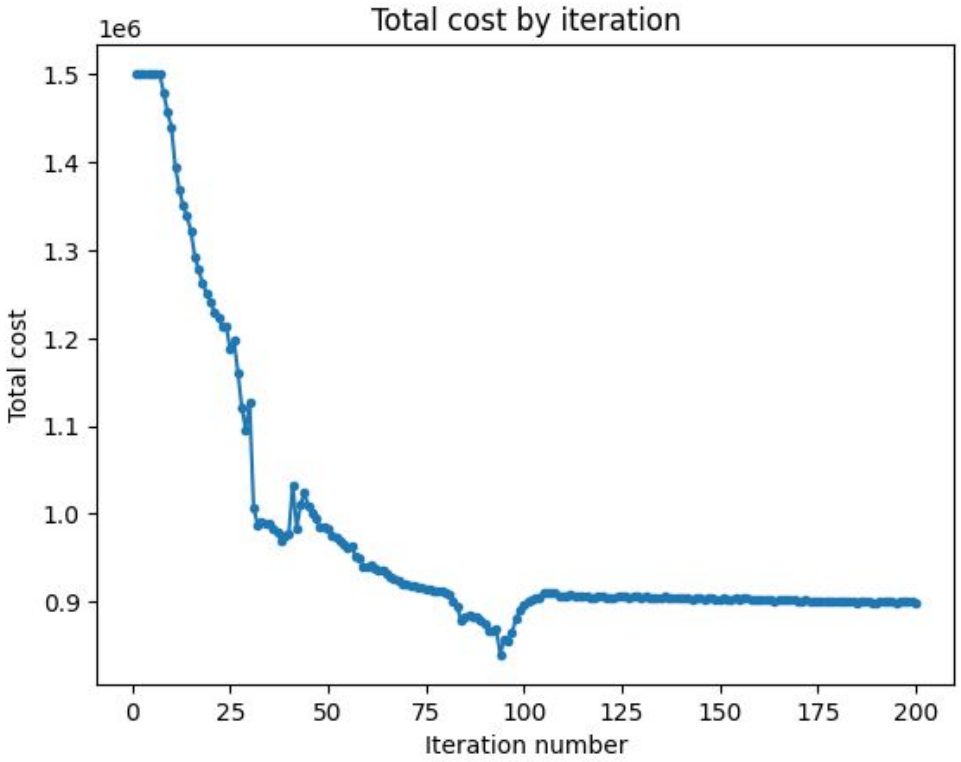

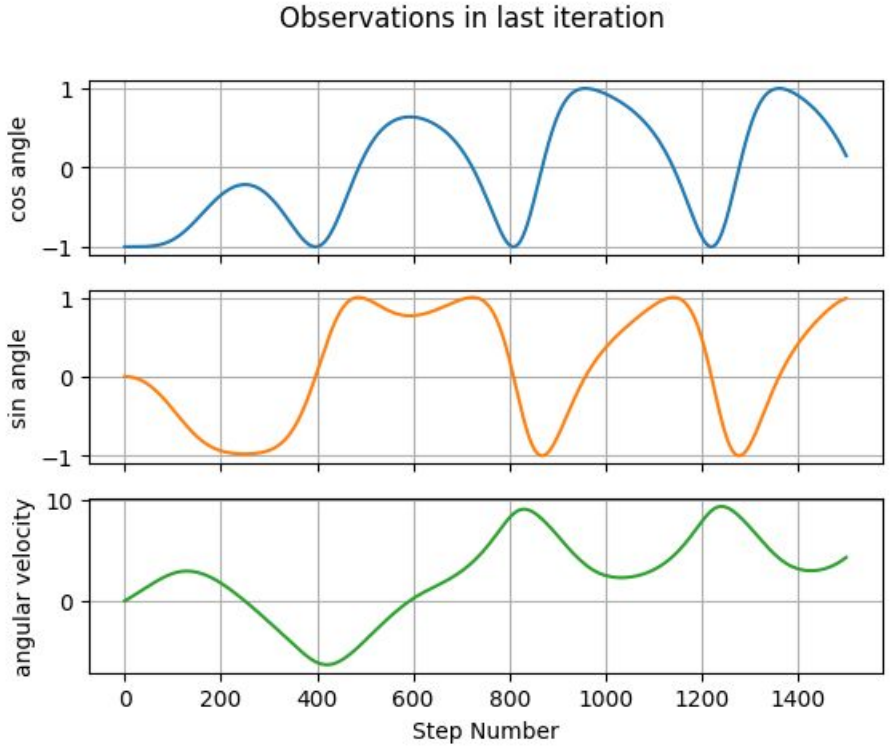

Hidden dim: 32
Hidden layers: 2

lr: 0.005
steps: 1500

N_episode: 2
N_iters: 200

cost function is fixed with velocity cost if angle is small (cosine > 0.95)

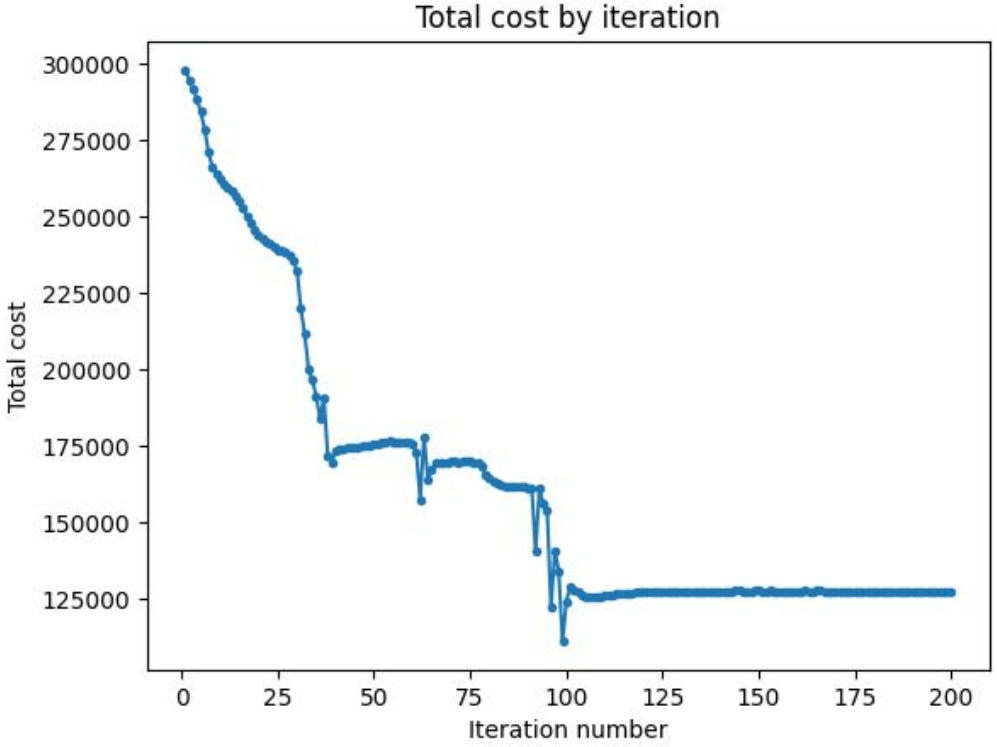

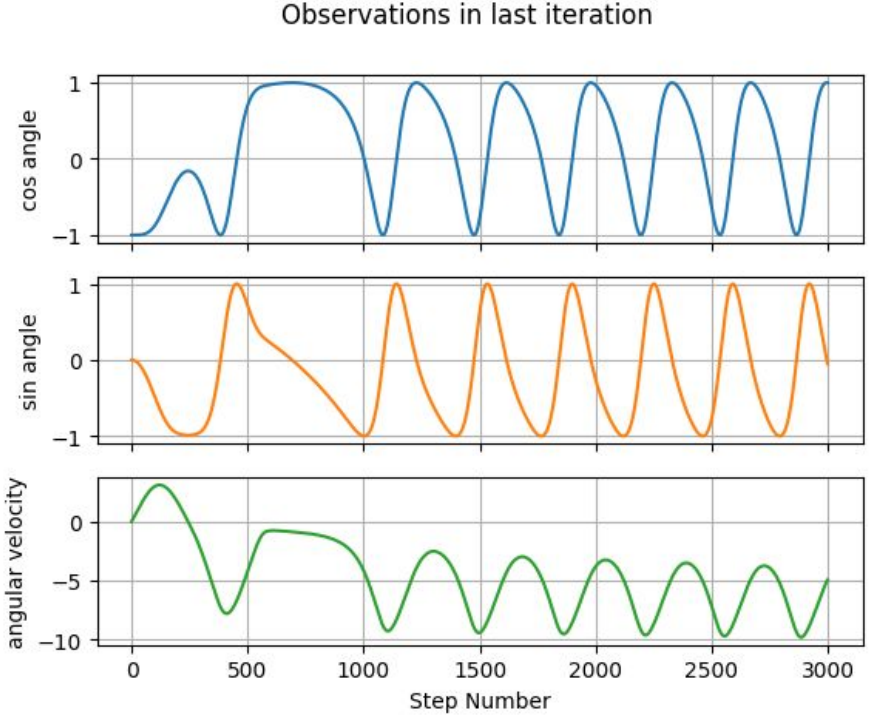

Hidden dim: 32 lr: 0.005 N_episode: 2
Hidden layers: 2 steps: 1500 N_iters: 200

cost function is $50 * (1 - \cos\theta)$ with velocity cost if angle is small (cosine > 0.95)

Setup 4

Actor Critic approach. Unlimited surface of cart.

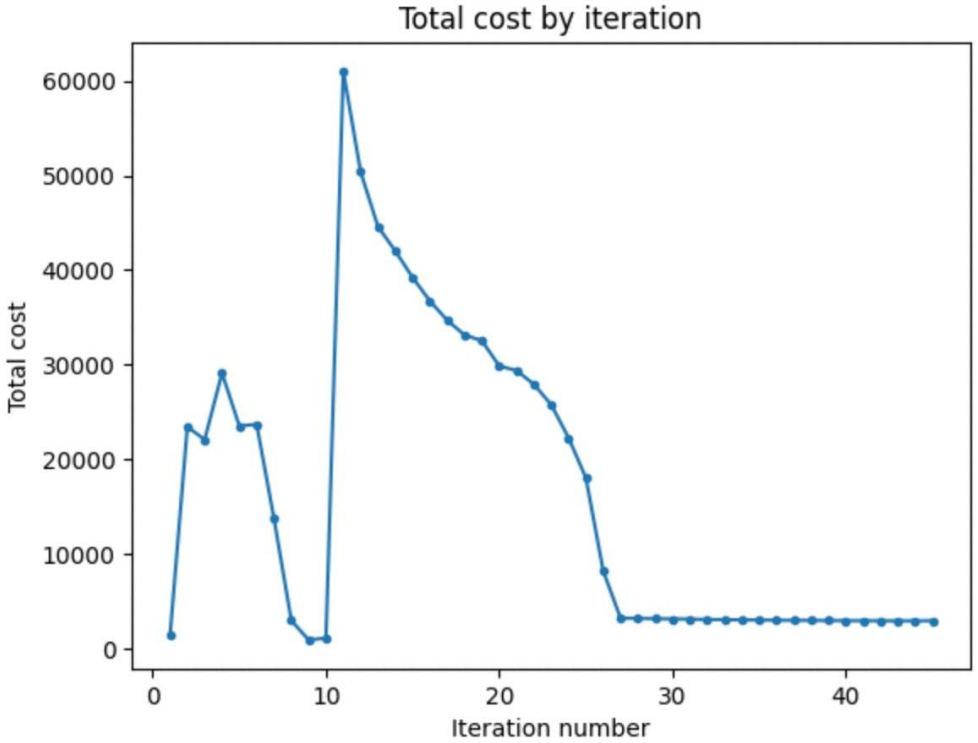

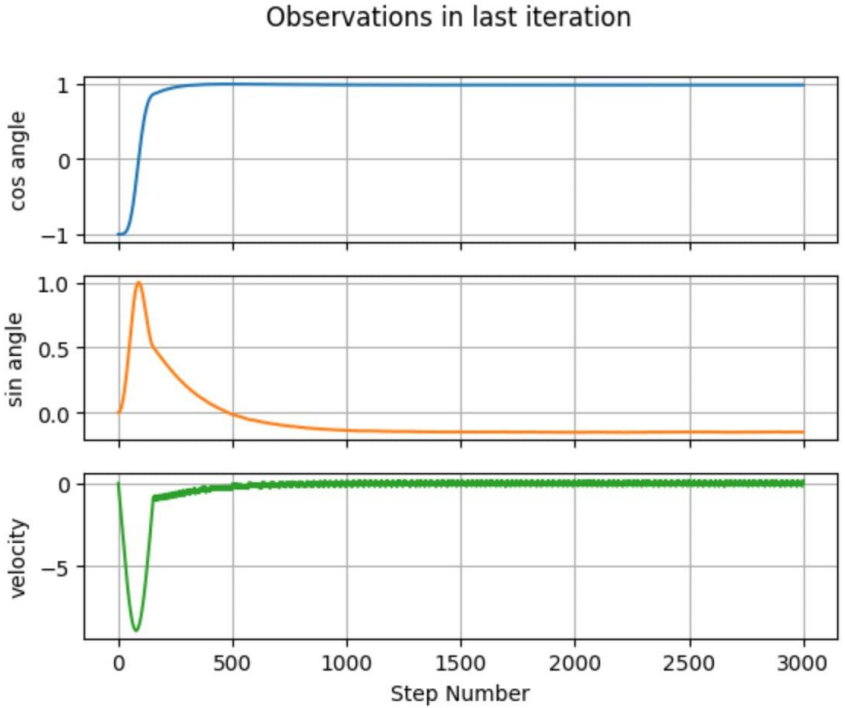

Hidden dim: 64 (both A-C)
Hidden layers: 8 (both A-C)
bounds: [-40, 40]
cost function is $(1 - \cos\theta) + \omega^2$

lr: 0.005
steps: 1500

N_episode: 2
N_iters: 200

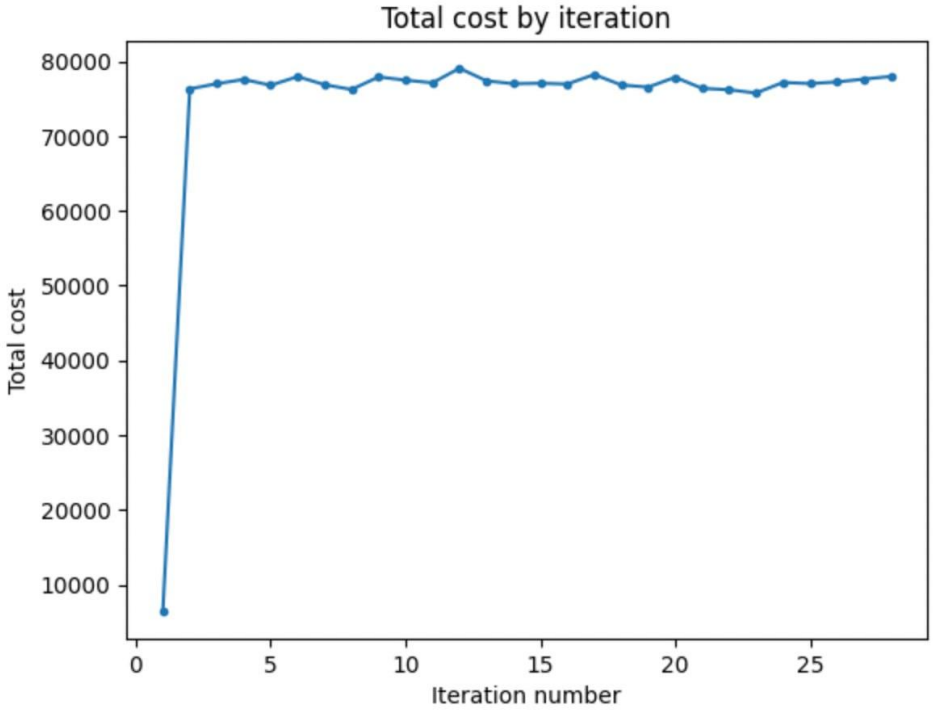

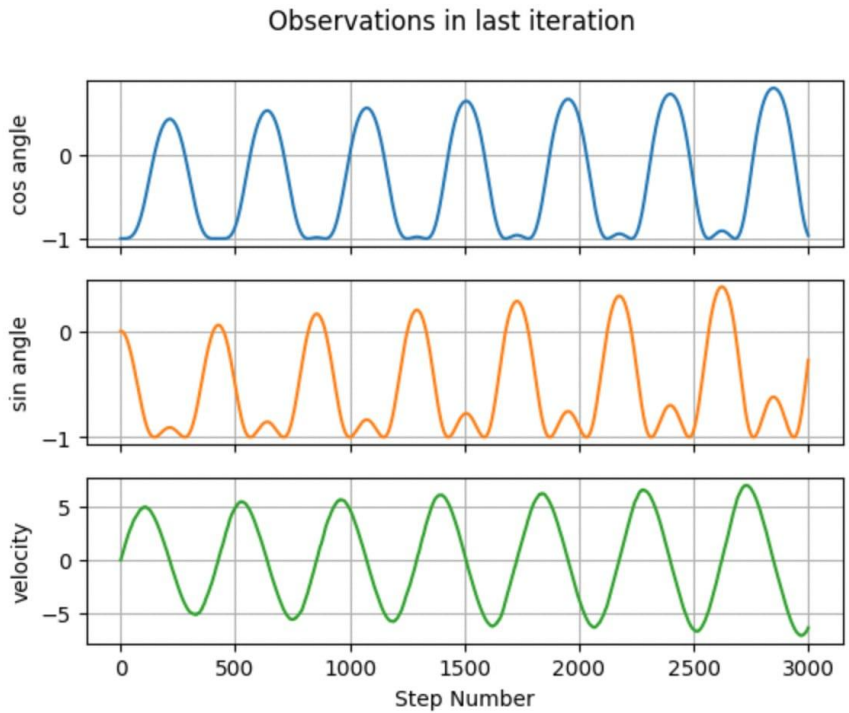

Hidden dim: 64 (both A-C) lr: 0.005 N_episode: 2
Hidden layers: 8 (both A-C) steps: 1500 N_iters: 200
bounds: [-40, 40] alpha: 0.5
cost function is $alpha * (1 - \cos\theta) + (1 - alpha) * \omega^2 + 0.25 * u^2$

Limited Cart Pole (-10, 10)

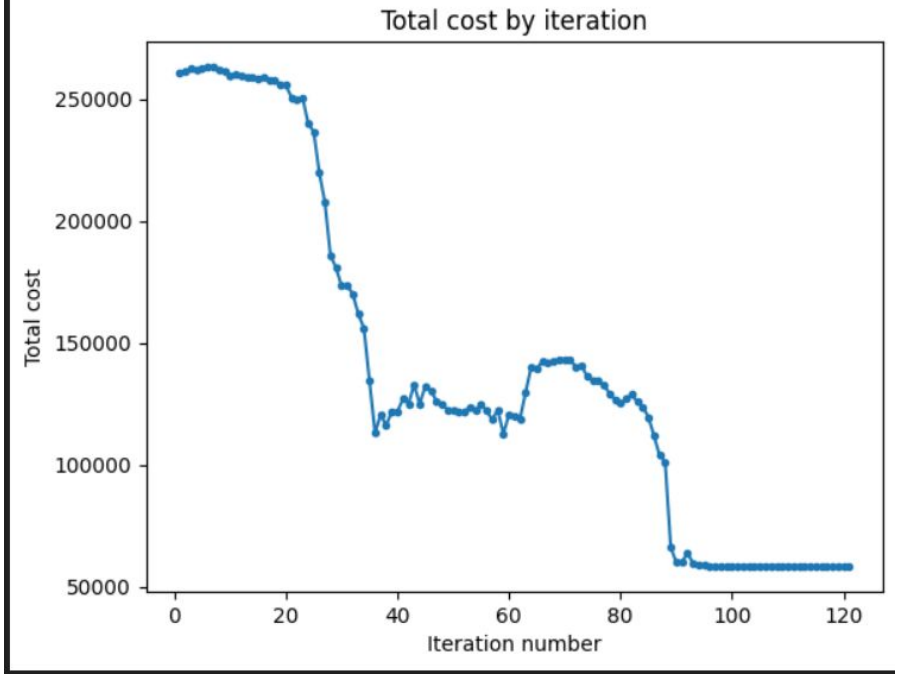

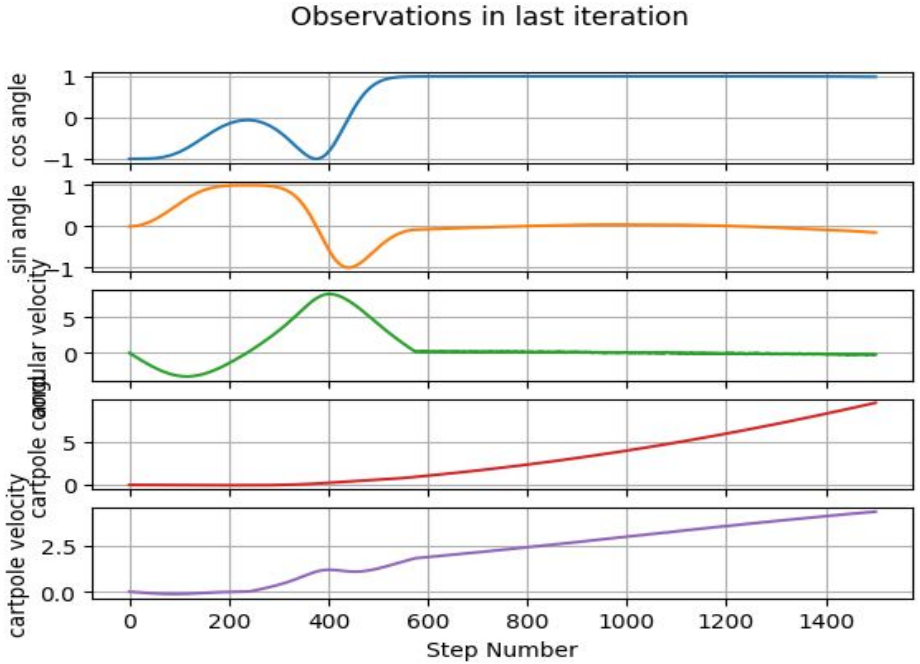

Hidden dim: 64
Hidden layers: 3

lr: 0.005
steps: 1500

N_episode: 2
N_iters: 200

if $\cos\theta$ < 0.9: return $80 * (1 - \cos\theta) + \omega^2 + x^2$
else: return $40 * (1 - \cos\theta) + \omega^2$

As you may observe, RL suffers with stabilizing system at desired position (angle == 0).

Let’s check, if initial state is already at (0, 0) for unlimited cart pole. And compare with LQR.

LQR formula

$$\dot{x} = \begin{bmatrix} 0 & 1 \\ \frac{(M+m)g}{LM} & 0 \end{bmatrix} x + \begin{bmatrix} 0 \\ -\frac{1}{LM} \end{bmatrix} u$$

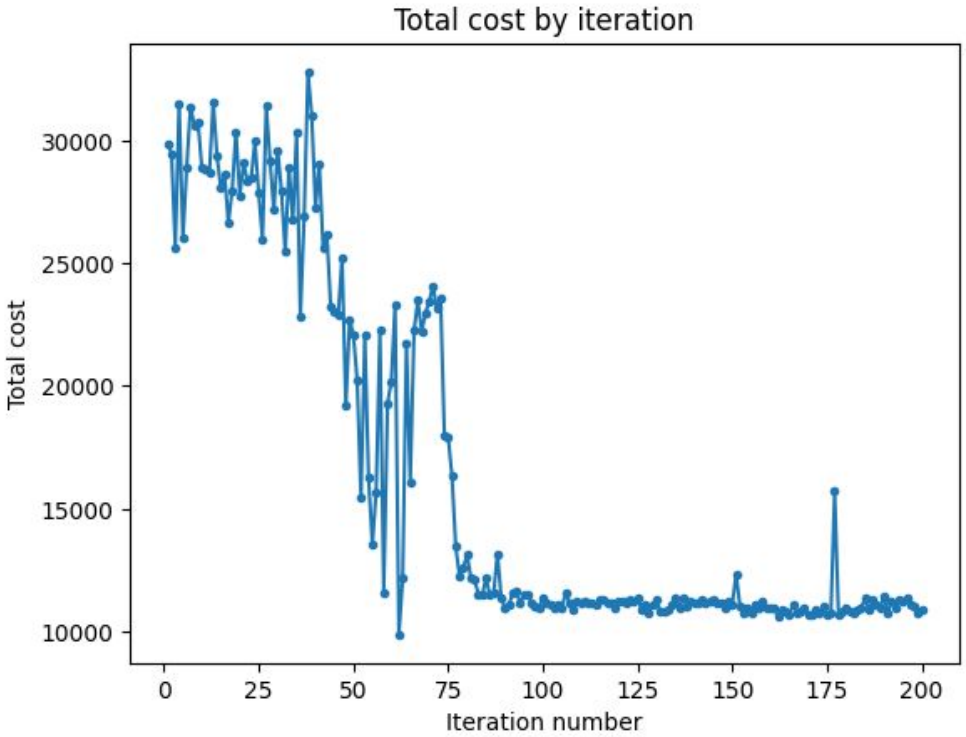

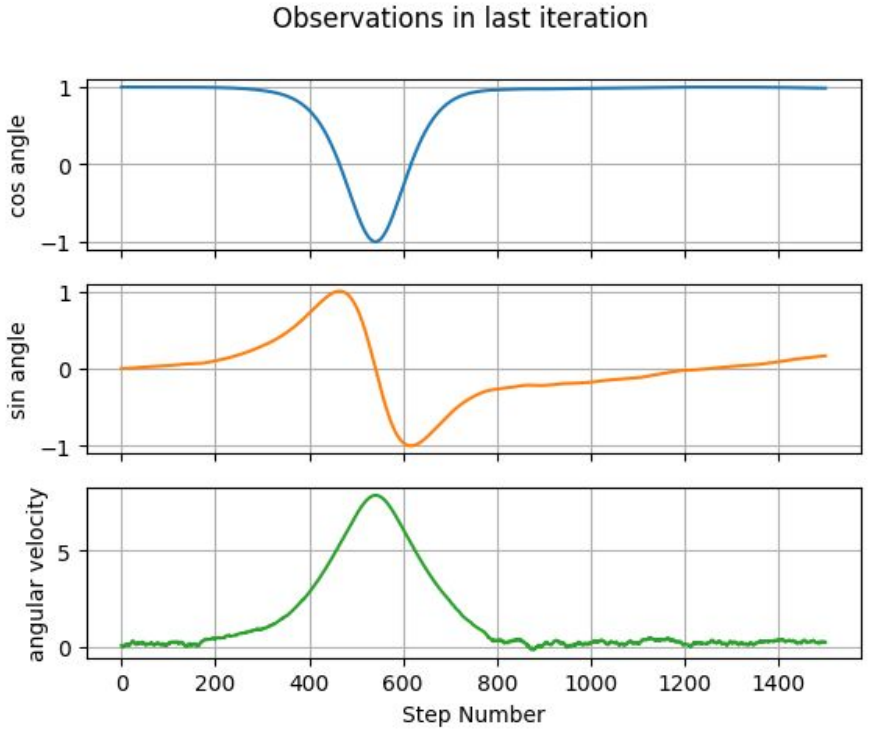

Hidden dim: 32
Hidden layers: 2

lr: 0.005
steps: 1800

N_episode: 2
N_iters: 200

cost fixed with velocity punishment on high cosine.

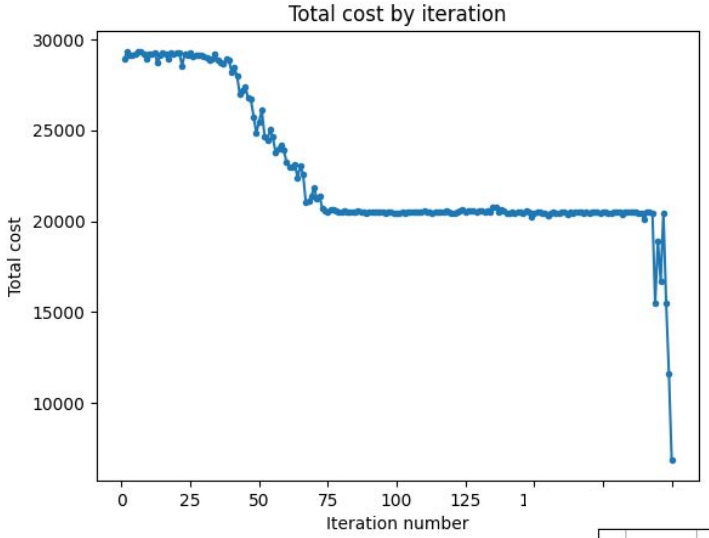
Total cost by iteration
Total cost
Iteration number
30000
25000
20000
15000
10000
0
25
50
75
100
125

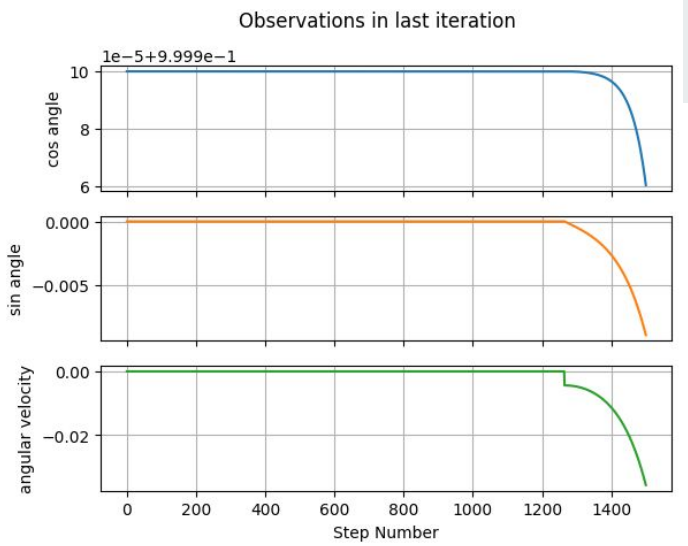
Observations in last iteration
1e−5+9.999e−1
cos angle
sin angle
angular velocity
Step Number
10
8
6
0.000
−0.005
0.00
−0.02
0
200
400
600
800
1000
1200
1400

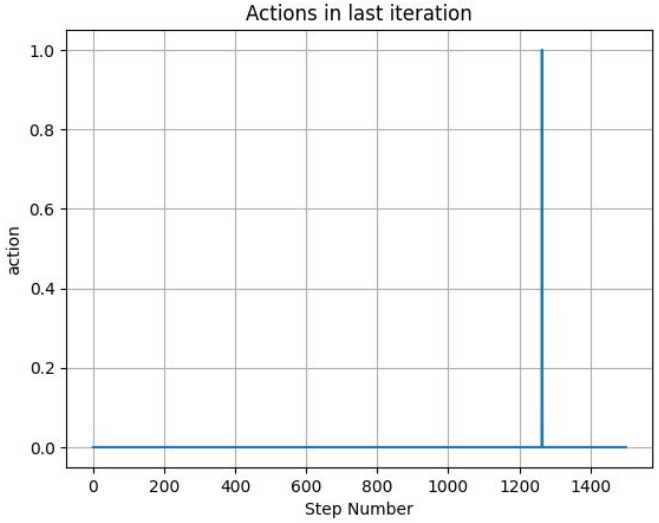
Actions in last iteration
action
Step Number
1.0
0.8
0.6
0.4
0.2
0.0
0
200
400
600
800
1000
1200
1400

Init state (np.pi / 100, 0)

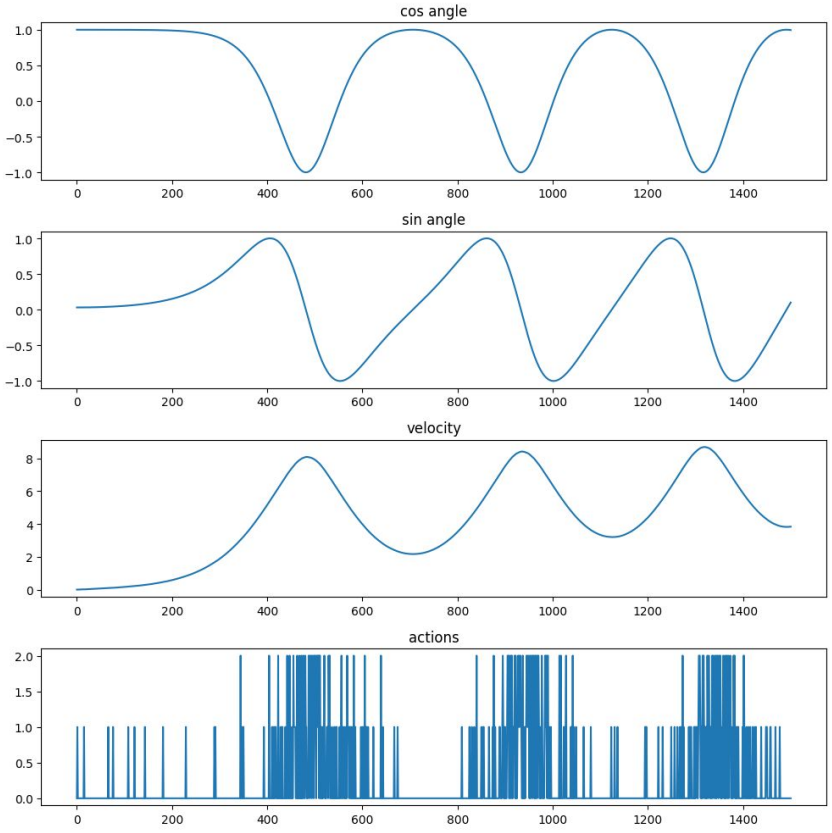

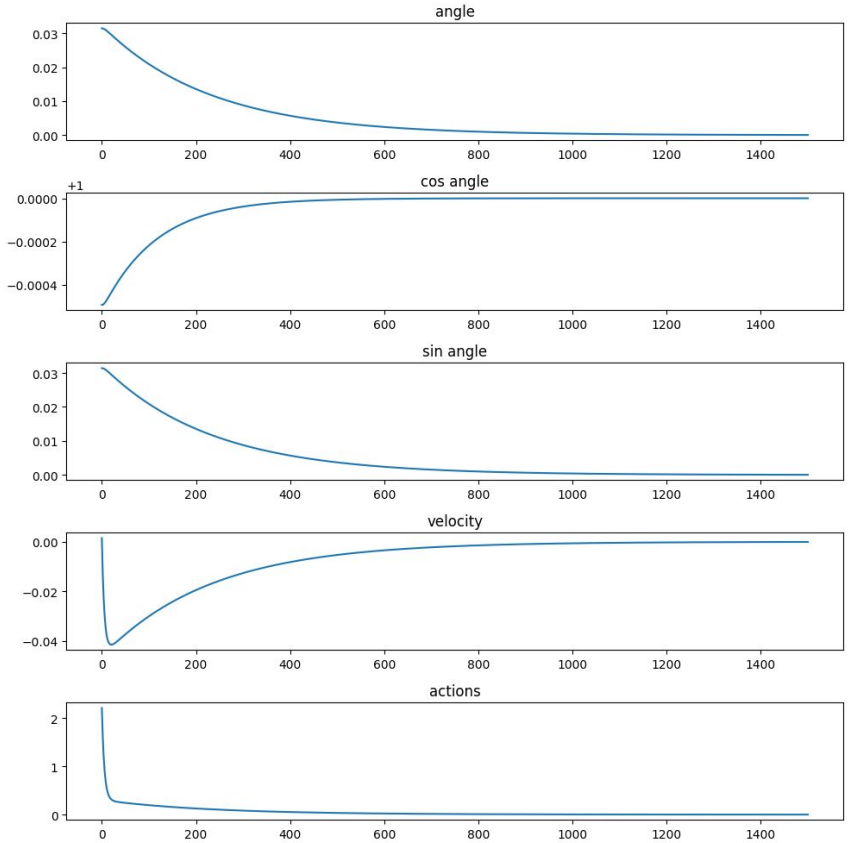

Thank you for your attention!

Ready to answer your questions.

https://github.com/BogChamp/rl_project/tree/master